

Department of Electronical Engineering, Telecommunications and
Computers

Remote Lab

50565: Ângelo Azevedo (a50565@alunos.isel.pt)
50539: António Alves (a50539@alunos.isel.pt)

Report for Project and Seminar Class
of Computer Science and Computer Engineering BSc

Advisor: Prof. Pedro Miguens Matutino

LISBON SCHOOL OF ENGINEERING

Remote Lab

50565 Ângelo Azevedo

50539 António Alves

Advisor: Pedro Matutino

Report for Project and Seminar Class of Computer Science and Computer Engineering
BSc

June 2025

Abstract

The design, development, implementation, and validation of digital systems require, in addition to simulators, the use of hardware for verification of their implementation in real devices. However, access to these real devices is sometimes restricted, not being available 24/7. In the current teaching paradigm where face-to-face time is reduced and remote and autonomous work is increased, it is necessary to create alternatives to the usual model.

The Remote Lab project aims to provide a virtual lab with access to remote hardware. This lab consists of a web application running on an embedded system. The web application, accessed via a website, aims to provide a dashboard where users can join a laboratory. This is where users can control the remote hardware. A hierarchy system will be implemented to provide different roles, each with their own permissions relative to how users can browse the information provided by the web application.

This project will implement the infrastructure to support the configuration, manipulation and visualization of remote hardware. Based on an architecture with back-end (database and Web API) and front-end (Web App, with a dashboard).

Resumo

A concepção, desenvolvimento, implementação, e por fim a validação de sistemas digitais requerem para além dos simuladores, a utilização de hardware para uma verificação da sua concretização em dispositivos reais. No entanto, o acesso a esses dispositivos reais é por vezes restrito, não estando acessíveis 24h/7. No atual paradigma de ensino em que se reduz o tempo presencial, aumentando o trabalho remoto e autónomo, é necessário criar alternativas ao modelo habitual.

O projeto Remote Lab tem como objetivo fornecer um laboratório virtual com acesso a hardware remoto. Este laboratório consiste numa aplicação web executada num sistema embebido. A aplicação web, acedida através de um website, visa fornecer um dashboard onde os utilizadores podem aderir a um laboratório. É aqui que os utilizadores podem controlar o hardware remoto. Será implementado um sistema hierárquico para fornecer diferentes funções, cada uma com as suas próprias permissões relativamente à forma como os utilizadores podem navegar pela informação fornecida pela aplicação web.

Este projeto implementará a infraestrutura de suporte à configuração, manipulação e visualização de hardware remoto. Baseado numa arquitetura com back-end (base de dados e Web API) e front-end (Web App, com um dashboard).

Contents

List of Figures	xi
List of Listings	xiii
Acronyms	1
1 Introduction	3
1.1 Context and Motivation	3
1.2 Objectives	3
1.3 Scope	3
1.4 Methodology	3
1.5 Structure of the Document	4
2 Placement	5
2.1 Related Work	5
2.2 Similar Systems	5
3 Proposed Architecture	7
3.1 System Overview	7
3.2 Main Components	7
3.3 Component Interactions	8
3.4 Architecture Diagram	8
3.5 Design Rationale	8
4 Implemented Infrastructure	9
4.1 Overview	9
4.2 Project Structure	10
4.3 Implementation Details	10
4.4 Database	10
4.4.1 Entity-Relationship Model	10
4.4.2 Implementation Details	14
4.4.3 Conclusion	14

4.5	Web API	14
4.6	Web Application	14
4.7	Deployment	14
4.8	Technologies Used	14
4.9	System Components	15
4.10	Deployment Architecture	15
4.11	Build and CI/CD	16
4.12	Notable Implementation Details	16
4.13	Summary	16
References		17

List of Figures

3.1	High-level architecture of the Remote Lab platform.	8
4.1	System Architecture Overview	9
4.2	Entity-Relationship Model (ER Model)	10
4.3	User Entity	11
4.4	Token Entity	11
4.5	Laboratory Entity	12
4.6	Hardware Entity	12
4.7	Group Entity	13
4.8	Lab Session Entity	13

Listings

Acronyms

API Application Programming Interface

BSc Bachelor of Science

ER Model Entity-Relationship Model

Chapter 1

Introduction

1.1 Context and Motivation

In the recent years, the need for remote access to laboratory resources has grown significantly, driven by the expansion of online education, research collaboration, and the increasing complexity of experimental setups. Traditional laboratories often require physical presence, which can limit accessibility and flexibility for students, researchers, and professionals. The **Remote Lab** project aims to address these challenges by providing a platform that enables secure, efficient, and user-friendly remote access to laboratory equipment and resources.

1.2 Objectives

The main objectives of the Remote Lab project are:

- To design and implement a scalable platform for remote laboratory access.
- To ensure secure authentication and authorization for different user roles.
- To provide an intuitive user interface for managing and scheduling laboratory sessions.
- To support integration with various types of laboratory hardware.

1.3 Scope

This project focuses on the development of the core platform, including backend services, user management, and basic hardware integration. Advanced features such as real-time data analytics, support for a wide range of laboratory devices, and extensive reporting capabilities are considered out of scope for the current phase.

1.4 Methodology

The project follows a modular and iterative development approach, leveraging modern software engineering practices. The backend is implemented using Kotlin and follows a layered architec-

ture, while the frontend is developed with Next.js to provide a responsive and accessible user experience.

1.5 Structure of the Document

The remainder of this report is organized as follows:

- **Chapter 2:** Related Work – Overview of existing solutions and technologies.
- **Chapter 3:** System Architecture – Description of the overall system design.
- **Chapter 4:** Implementation – Details of the main components and their interactions.
- **Chapter 5:** Evaluation – Assessment of the system’s performance and usability.
- **Chapter 6:** Conclusions and Future Work – Summary of achievements and directions for future development.

Chapter 2

Placement

This chapter is organized into two sections, where we describe related work and some systems similar to the one developed in this project.

2.1 Related Work

In recent years, several initiatives have emerged to provide remote access to laboratory resources, especially in the context of higher education and research. Projects such as MIT's iLab and LabShare have demonstrated the feasibility and benefits of remote laboratories, enabling students and researchers to conduct experiments from anywhere in the world. These platforms typically focus on providing secure access, scheduling, and integration with a variety of laboratory equipment. The literature highlights the importance of usability, scalability, and security in the design of such systems, as well as the challenges associated with real-time interaction and hardware integration.

2.2 Similar Systems

There are several systems that offer functionalities similar to those of the Remote Lab project. For example, the iLab Shared Architecture (ISA) provides a framework for sharing laboratory equipment over the internet, supporting both batch and interactive experiments. LabShare is another notable example, offering a collaborative platform for remote experimentation and resource sharing among institutions. Other systems, such as WebLab-Deusto and VISIR, focus on specific domains like electronics and instrumentation, providing specialized interfaces and tools for remote experimentation. These systems serve as valuable references for the development of the Remote Lab platform, informing decisions related to architecture, user experience, and integration with laboratory hardware.

Chapter 3

Proposed Architecture

This chapter presents the proposed architecture for the Remote Lab platform, detailing its main components, their interactions, and the rationale behind the architectural choices.

3.1 System Overview

The Remote Lab platform is designed as a modular and scalable system, enabling secure and efficient remote access to laboratory equipment. The architecture follows a layered approach, separating concerns between the user interface, application logic, and hardware integration. This separation facilitates maintainability, extensibility, and the integration of new features or laboratory devices.

3.2 Main Components

The architecture consists of the following main components:

- **Frontend:** A web-based user interface developed with Next.js, providing users with access to laboratory resources, session scheduling, and experiment monitoring.
- **Backend:** Implemented in Kotlin, the backend exposes RESTful APIs for user management, authentication, authorization, and laboratory session control. It also handles business logic and enforces security policies.
- **Hardware Abstraction Layer:** This layer manages communication with laboratory equipment, abstracting hardware-specific details and providing a unified interface for the backend.
- **Database:** Stores user data, session information, access logs, and configuration settings. The database ensures data consistency and supports auditing requirements.
- **Authentication and Authorization:** Ensures secure access to the platform, supporting multiple user roles (e.g., students, professors, administrators) with different permissions.

3.3 Component Interactions

The components interact as follows:

- Users interact with the frontend to authenticate, schedule sessions, and access laboratory resources.
- The frontend communicates with the backend via secure API calls.
- The backend processes requests, applies business logic, and interacts with the database and hardware abstraction layer as needed.
- The hardware abstraction layer translates backend commands into device-specific instructions, enabling remote control of laboratory equipment.

3.4 Architecture Diagram

Figure 3.1 illustrates the high-level architecture of the Remote Lab platform.

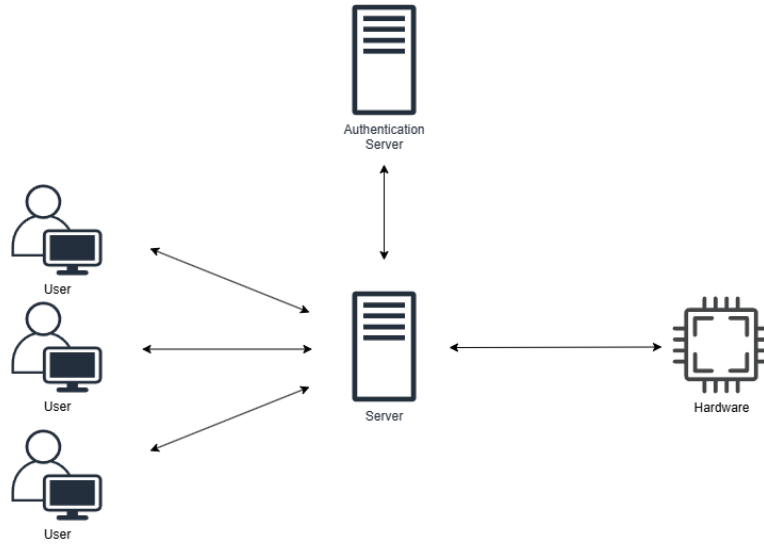


Figure 3.1: High-level architecture of the Remote Lab platform.

3.5 Design Rationale

The architectural choices were guided by the need for scalability, security, and ease of integration with diverse laboratory equipment. The use of a layered architecture and standardized interfaces ensures that the platform can evolve to meet future requirements and support additional functionalities.

Chapter 4

Implemented Infrastructure

This chapter details the infrastructure implemented for the Remote Lab platform, covering the main components, technologies, deployment strategy, and integration between system modules.

4.1 Overview

The Remote Lab platform is designed as a modular, containerized system that enables secure and efficient remote access to laboratory equipment. The infrastructure follows a layered architecture, separating the user interface, backend logic, and hardware integration, and is built with scalability and maintainability in mind.

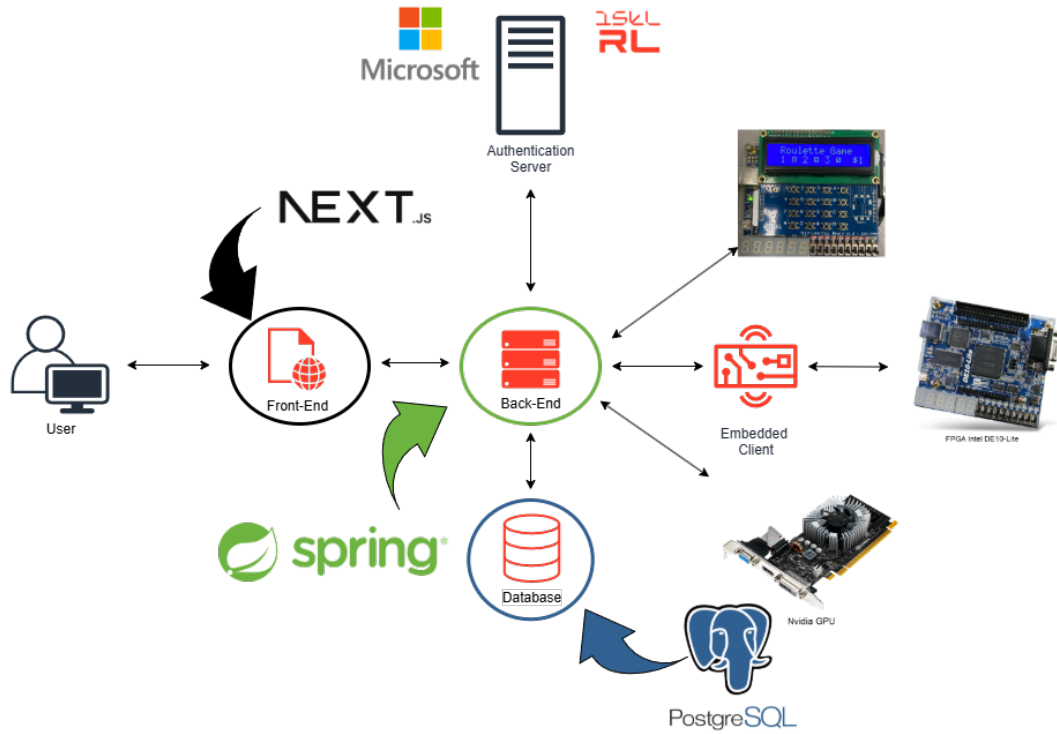


Figure 4.1: System Architecture Overview

4.2 Project Structure

Project structure explanation.

4.3 Implementation Details

Project decisions.

4.4 Database

The database serves as the foundational component of the system architecture. PostgreSQL was selected as the database management system due to its reliability, open-source nature, and robust support for relational data models. This choice aligns with previous project implementations and provides the consistency and performance required for the system's operational needs.

This section presents an overview of the Entity-Relationship Model (ER Model) and critical implementation details. Complete technical documentation is provided in the accompanying appendix.

4.4.1 Entity-Relationship Model

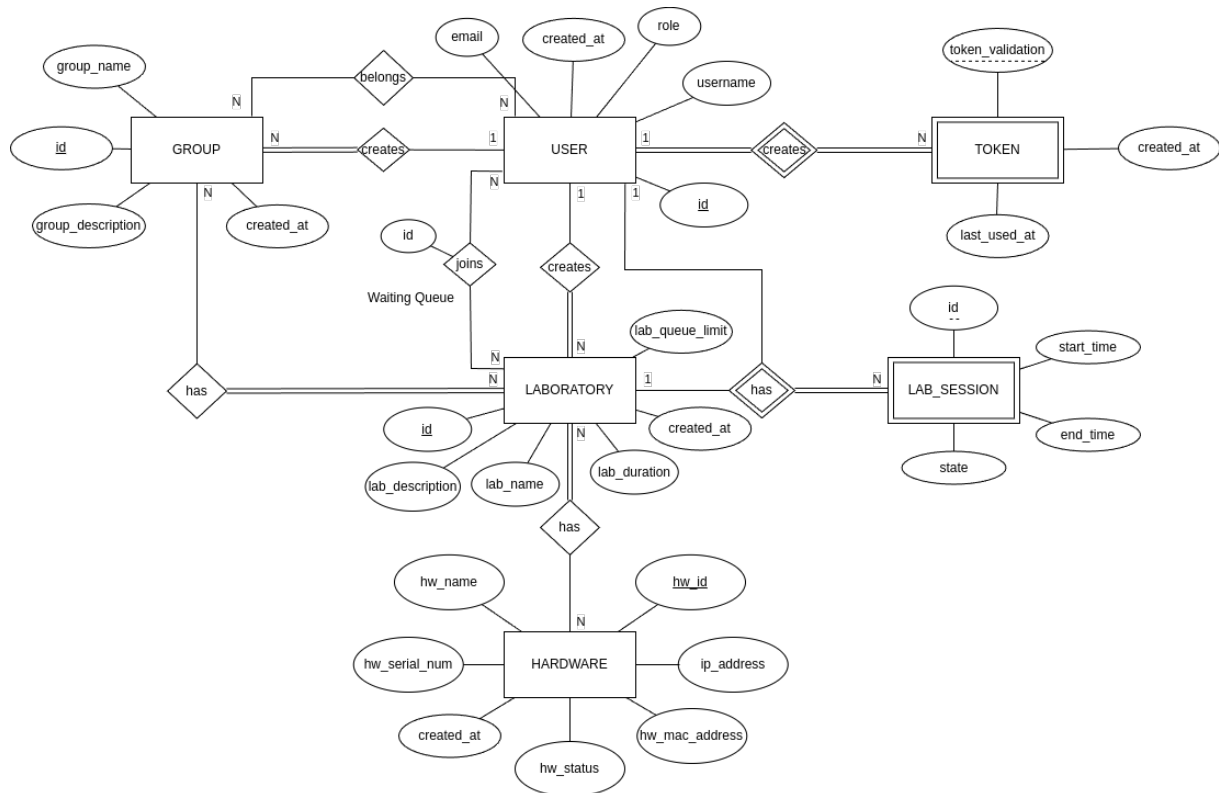


Figure 4.2: Entity-Relationship Model (ER Model)

The database design follows a normalized relational structure that supports user authentication, secure session management, and the remaining system functionalities. The ER model encompasses the core entities required for system functionality while maintaining data integrity and scalability.

Core Entities

User




Figure 4.3: User Entity

The **User** entity represents a user in the system. The username and email attributes are provided by the authentication system. The role serves as discriminator attribute to identify whether the user is an administrator, professor or student.

Token




Figure 4.4: Token Entity

A user can create N tokens. The **Token** is a weak entity because it cannot be identified by its attributes alone and therefore requires a user, which is a strong entity, to be identified. Its attributes cannot uniquely identify it. A token is created by only one user.

This is a useful entity for authentication purposes. It was designed to hold a hash value in the *token_validation* attribute.

Authentication workflow:

1. Upon successful user login, a unique token is created with cryptographically secure values and stored in the database.
2. For subsequent authenticated operations, the system queries the database to verify the client-provided token against stored values.
3. Valid tokens enable secure user identification without transmitting unique identifiers.

The *last_used_at* and the *created_at* are useful for determining token expiration.

Laboratory




Figure 4.5: Laboratory Entity

A user, as an administrator or professor, can create N laboratories. When creating a **Lab-
oratory**, the user can define the name (*lab_name*) and description (*lab_description*). They can also define the duration of a laboratory session (*lab_duration*) and its queue limit (*lab_queue_limit*).

Hardware




Figure 4.6: Hardware Entity

Upon successful laboratory creation, the user can associate **Hardware** to it, which must be created separately.

For the creation, it requires a name (*hw_name*), IP (*ip_address*) and MAC (*mac_address*) addresses (which can be null depending on the hardware), a status (*hw_status*) to indicate whether the hardware is under maintenance, occupied, or available, and a serial number (*hw_serial_num*) to uniquely identify the hardware. Although it has an ID, the serial number helps physically identify the hardware.

Group




Figure 4.7: Group Entity

For a student to access a laboratory, they must be in a group that is associated with that laboratory. A professor can create a **Group** and associate users to it.

When creating a group, the user needs to name it (*group_name*) and, optionally, add a description (*group_description*) to it.

Lab Session




Figure 4.8: Lab Session Entity

Finally, a user can join a laboratory if they are in a group associated with it. If the laboratory is being used, the user enters a waiting queue; otherwise, a **Lab Session** is created.

Lab Session is a weak entity. It requires two strong entities to be identified: the **User** entity and the **Laboratory** entity. This is used to check whether a user is in a lab session or for statistical purposes. The *state* attribute indicates whether the session is over or still running. The *start_time* and *end_time* can be used for statistical details, such as determining how much time a user spent in a laboratory, or for future purposes, such as scheduling sessions.

4.4.2 Implementation Details

After providing an overview of the database entities and their associations, there are important details worth mentioning:

- Although PostgreSQL is being used for its functionalities, it was decided that all logic and verifications are implemented in the Web API, so that no triggers or complex constraints are implemented on the database side.

4.4.3 Conclusion

This section has provided an overview of the database architecture, implementation, and design decisions. It has also presented the ER Model of the database and described a typical user journey, explaining database interactions.

The documentation should be consulted for a comprehensive deep dive. It explains every entity, its attributes, and provides theoretical insights.

4.5 Web API

Web API

4.6 Web Application

Web App

4.7 Deployment

4.8 Technologies Used

- **Frontend:** Implemented with Next.js (React framework), providing a modern, responsive web interface for users to interact with laboratories, schedule sessions, and control hardware.
- **Backend:** Developed in Kotlin using Spring Boot, exposing RESTful APIs for user management, authentication, laboratory session control, and business logic enforcement.

- **Database:** PostgreSQL is used to persist user data, session information, access logs, and configuration settings.
- **ORM/Database Access:** JDBI is used for type-safe, modular database access in the backend.
- **Authentication:** Microsoft OAuth via NextAuth is used for user authentication, supporting multiple roles (student, professor, administrator).
- **Containerization:** Docker is used to containerize all major components (frontend, backend, database), ensuring consistent deployment across environments.
- **Orchestration:** Docker Compose manages multi-container deployment, networking, and environment configuration.

4.9 System Components

- **Web Application (Frontend):** Provides dashboards, laboratory access, real-time hardware monitoring, and session management. Built with Next.js and deployed as a Docker container.
- **API Server (Backend):** Handles authentication, authorization, laboratory and user management, and hardware abstraction. Built with Kotlin and Spring Boot, also containerized.
- **Database:** PostgreSQL instance running in a Docker container, with persistent storage volumes.
- **Hardware Abstraction Layer:** Backend modules abstract hardware-specific details, exposing unified interfaces for laboratory equipment control.

4.10 Deployment Architecture

The system is deployed using Docker Compose, which defines and manages the following services:

- **db:** PostgreSQL database container, with health checks and persistent volumes.
- **api:** Backend API container, built from the Kotlin/Spring Boot project, depending on the database service.
- **website:** Frontend container, built from the Next.js project, depending on the API service.

All services are connected via Docker networks to ensure secure and efficient communication. Environment variables and secrets are managed via `.env` files.

4.11 Build and CI/CD

- **Gradle:** Used for building and managing backend dependencies.
- **NPM:** Used for frontend dependency management and builds.
- **Dockerfiles:** Multi-stage builds are used for both backend and frontend to optimize image size and security.
- **GitHub Actions:** (If applicable) Used for continuous integration and automated builds.

4.12 Notable Implementation Details

- The backend uses JDBI for database access, configured with application-specific requirements.
- Environment variables are used to configure database connections and secrets, improving security and flexibility.
- The system supports role-based access control, with different permissions for students, professors, and administrators.
- The hardware abstraction layer allows for future extension to new types of laboratory equipment.

4.13 Summary

The implemented infrastructure leverages modern web technologies, containerization, and modular design to provide a robust, scalable, and maintainable platform for remote laboratory access.

References